

* Blocking Factor [Bf]: Avg no of records / block.

$$\rightarrow Bf = \frac{\text{Block Size}}{\text{Record Size}} \quad (\text{Block size} > \text{Record size})$$

Spanned Organisation

(1) A record can't be stored in more than one block.

$$(2) Bf = \frac{\text{Block size}}{\text{Record size}}$$

(3) No memory waste, but block access cost increased.

(4) Suitable for variable length records.

Unspanned Organisations

Record can be stored in a particular block.

$$Bf = \left\lceil \frac{\text{Block size}}{\text{Record size}} \right\rceil$$

Block access cost reduced but waste of memory.

Suitable for fixed length records.

* Default is unspanned organisation.

Block factor: max possible records / block.

→ for unspanned:

$$\left\lceil \frac{B+H}{R} \right\rceil \text{ records / block}$$

H = header size.

→ for spanned:

$$= \frac{B+H}{R}$$

Ordered file

1. Searching easy.
2. Insertion expensive.
3. To access any record
blocks = $\lceil \log_2 B \rceil$

Unordered File

- Insertion easy.
Searching expensive.
- Linear Search:
- $$\text{Avg case} = \frac{B}{2}$$
- $$\text{Worst case} = B$$

One Index Recordfile

= Size of search key +
Size of pointer.

$$\# \text{data blocks} = \frac{\# \text{records}}{\text{Bf of DB file}}$$

Dense Index

* Index entry created for every search key value.

* No. of index entries = no. of DB records.

Sparse Index

Index entry created for some search key values.

No. of index entries = no. of blocks.

Bf of Index file

$$\left[\frac{\text{Block size}}{\text{Index record size}} \right]$$

$$\rightarrow \text{Block Factor of Index} = \left[\frac{B-1}{k_{sp}} \right] \text{ entries/blocks}$$

Single level Ordered Index:

1. Primary Index (key + ordered file)

2. Clustered Index (nonkey + ordered file)

3. Secondary Index (key/nonkey + order/unordered file)

* Atmost one primary index / relation, and 21 secondary index possible.

* In a relation either CI or PI but not both

\rightarrow Access cost, PI using multi-level index = (k_1) blocks

\rightarrow To access a record using index acc no of block = $\lceil \log_2 \text{Blocks} \rceil + 1$

\rightarrow 'n' level in multi-level index, # blocks acc to search a record

$$= n+1$$

B-Tree

P: block/child/index pointer

Order Φ : max possible child

Pr: read/record/data pointer

pointers that can be stored

k: Search key

$$\underline{\text{more}} - B_p = P, \underline{\text{more keys}} (P-1)$$

\rightarrow Every internal node: $B_p = \lceil \frac{P}{2} \rceil$, min keys $(\frac{P}{2}-1)$
(except root)

\rightarrow Root can contain at least $-2B_p$ (min 1 key)

\rightarrow All leaf nodes are same level. Each leaf node should bear same level.

\rightarrow Order: P by default/min (P-block pointer)

$$(T_{max})_{\text{node}} = P + B_p + (p_1) (\text{key} + P) \leq \text{BlockSize}$$

→ Order: P (key) (key order: p)

$$= (P+1) B_p + P(\text{key} + R_p) \leq \text{Block Size}$$

Order p :	height/level	Min. node	Min #Bp	Min #keys
Minimum	$b/b+1$	$2\lceil \frac{P}{2} \rceil^{b+1}$	$2\lceil \frac{P}{2} \rceil^b$	$2\lceil \frac{P}{2} \rceil (\lceil \frac{P}{2} \rceil - 1)$

Order p :	height/level	Max. node	Max #Bp	Maximum keys
Maximum	$b/b+1$	P^b	P^{b+1}	$P^b(P-1)$

Total minimum keys: $(1 + 2(\lceil \frac{P}{2} \rceil - 1) + 2\lceil \frac{P}{2} \rceil (\lceil \frac{P}{2} \rceil - 1) + \dots + 2\lceil \frac{P}{2} \rceil^{b+1} (\lceil \frac{P}{2} \rceil - 1)) \rightarrow \underline{\text{GP formula}}$

Total maximum keys: $(P-1) + P(P-1) + P^2(P-1) \dots P^b(P-1) \rightarrow \underline{\text{GP series}}$

B^+ Tree.

→ All keys are available at leaf node

→ Internal node: no read pointer.

→ Leaf node contains key, and $R_p + 1$ Block pointer

Order of internal node: $P * B_p + (P-1) \text{ key} \leq \text{Block Size}$

Order of leaf node: $(P-1)(\text{key} + R_p) + 1 B_p \leq \text{Block Size}$

→ Split of internal nodes is same as in B tree, but split of leaf is

different from B tree

→ B^+ tree best suited for sequential access of range of records

- Procedural Query language: Relational Algebra
- Non procedural QL: SQL and TRC, and DRC

Relational Algebra

- * Query executed on DB tuple by tuple. One tuple at a time.
- * Generates distinct tuples.
- * Takes relational instance as I/p and returns relational instance as O/p

Basic Operators

1. Selection (σ)
2. Projection (π)
3. Cross product (\times)
4. Unions (\cup)
5. Set difference ($-$)
6. Rename (ρ)

Derived Operators

1. Join and its types (\bowtie)
2. Division (\mid)
3. Intersection (\cap)

* To apply Set Operators ($\cup, \cap, -$), relation must be Union compatible:

→ Any of two relations must be same

→ range (domain) of attribute must be similar

Cross Product (\times)

$R \times S = n_1 \times n_2$ tuples,

($C_1 + C_2$ attribute)

(For cross product, relations not required to be Union compatible)

→ Cross joins = Cross product (\times)

→ Natural Joins = Natural Cross product

— when join-condition is empty

— no common attribute

$$\rightarrow \frac{\pi_{AB}(R)}{\pi_A(S)} = B, \quad \frac{\pi_{AB}(R)}{\pi_B(S)} = A$$

Inner Joins

1. Conditional Joins (\bowtie_C)
2. Equi Joins
3. Natural Joins

Outer Joins

1. Left outer join (\bowtie_L or \bowtie_{RL})
2. Right outer join (\bowtie_R or \bowtie_{LR})
3. Full outer join (\bowtie_F or \bowtie_{LR})

Structured Query Language

- * DDL : Create, Alter, Drop table
- * DML : Insert, Update, Delete
- * DCL : Commit, Abort.
- * DQL : Select, From, Where

→ SQL is not case sensitive.

→ Mandatory SQL clause are "SELECT" and "FROM".

→ SQL works on Tuple Set, retains duplicates.

SQL Clauses

1. Select
2. From
3. Where
4. Group By
5. Having
6. Order By, etc.

Aggregate Functions

1. ~~Avg~~ COUNT
2. ~~Avg~~ SUM
3. ~~Avg~~ AVG
4. ~~Avg~~ MAX
5. ~~Avg~~ MIN

* Aggregate functions

Cannot be lowercase.

NULL : not existed / unknown

* Arithmetic Operations uses "NULL" gives NULL

* Aggregate functions ignore NULL.

* Comparison with NULL gives value 'Unknown'

Set Operations

1. Union & Union all
2. Intersect & intersect all.
3. minus & minus all
except except all

→ When aggregate Operator and other Attribute is used in Select clause, it is allowed only if attribute mentioned in Group by clause.

Regular Expressions

* % : Zero or more characters

* _ : Exactly one character.

A. → Starts with A

1. j → Ends with j

1. I. → Contains I.

_ → All lengths

'S---' → Starts with S, exactly 4 char.

'S---.' → Starts with S, min 4 char.

Tuple

Relational Calculus

Query described in the form of tuples.

Tuple variable:

Variable that takes on tuple of a relation schema as value.

Domain Relational Calculus

Query described result is the form of columns (Domain)

Domain variable

Variable that range over the domain of some attribute

Entity Set: Collection of similar entities that share the same properties.

Relationship: Associations among different entities.

Attribute: Properties used to describe an entity.

Attribute Types are:

1. Simple and Composite
2. Single and multivalued
3. Stored and derived
4. Key
5. Complex
6. Descriptive.

Mapping Cardinality Constraints

- * One to one ($\leftrightarrow \text{or } \square \rightarrow \square$)
- * One to many ($\leftrightarrow \text{or } \square \rightarrow \star$)
- * Many to one ($\square \rightarrow \square \text{ or } \square \rightarrow \star$)
- * Many to many ($\square \rightarrow \star \text{ or } \star \rightarrow \star$)

Referencing relation: Table using foreign key (Child table).

Referenced relations: Table referenced by foreign key (Parent table)

- Deletions from 'referenced' relations and insertions into 'referencing' relations may violate foreign key constraint.
- Foreign key may contain duplicates and null values.

ER to DBMS Conversion

ER	DBMS
(1) One : one	2 Tables req.
(2) 1 to M: many	2 Tables req.
(3) M : m	3 Tables req.
(4) M : 1	2 Tables req.

Partial participation at both sides
with above mapping are
Same as above.

* Partial Participation (\subset)

* Total / Full Participation ($=$)

→ If Total participation at both sides

In (1:m, m:1, 1:1, and m:m) then
1 table required

* If at one side is 1:1, then
Only 1 table required

* If at any side is A:m:m then
2 tables required

ACID Properties

1. Atomicity
2. Consistency
3. Isolation
4. Durability

Schedule: Sequence of instructions - that

Specify the chronological order in which

Instructions of current transaction are

executed

Serial schedule: Execution of one by one

Non Serial schedule: interleaved executions of

(may/may not be two or more transactions)

consistent) for 'n' transactions, n! serial schedules possible

Conflict Serializable: Transforming non serial to serial schedule by

Swapping

T_i T_j

Swapping non conflict instruction.

R(A) → W(A)

R(A) → R(A)

Conflicting instruction

W(A) → R(A)

R(A) → W(B)

→ i = read(Q), j = write(Q)

W(A) → W(A)

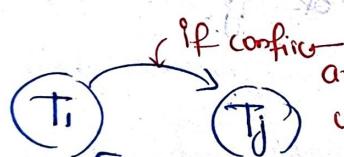
(Swapping)

→ i = write(Q), j = read(Q)

(Swapping not
possible)

Possible)

→ i = write(Q), j = write(Q)



arise,
we draw an
arc

→ If schedules S₁ and S₂
are conflict equal then
Procedure

Graph
Method

their precedence graph
will be same

(Ex) CNC

- A schedule is Conf. Serializable if it's either Conf. / View Serializable or both
- If a schedule is Conf. Serializable then it's already View Serializable, but vice versa not true.

View Serializable

- * If in schedule S, T_i reads initial value of Q, then in S also T_i must read the initial value of Q.
- * If in S, T_i executes read(Q), and that value was produced by T_j, then S in S also transactions must read the initial

- Q that was produced by same write (Q) operation of T_j.
- * The transaction that performs final write (Q) is S, must also perform final write (Q) in S.

View Equivalence of S₁ and S₂

- (1.) Initial reads of S₁ and S₂ should be same.
- (2.) Final updates of every data item should be same in S₁ and S₂.
- (3.) Write-read sequence should be equal.

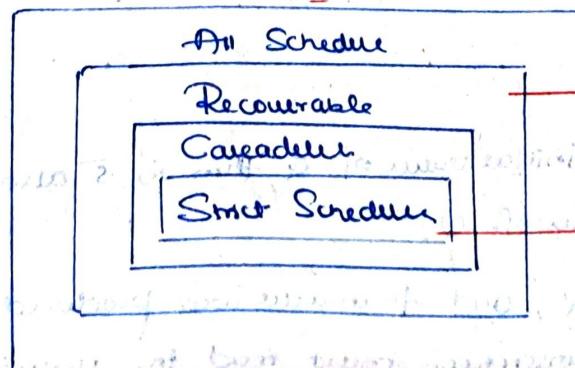
→ Non-serial Schedule = Total Concurrent - Serial Schedule (no!)

Recoverable Schedule: T_j reads a data item previously written by T_i then "Commit" of T_i appears before Commit of T_j.

Finding no. of Confice Serializable

1. Write down all operations of following (later) transactions.
2. Start from last operation of first transaction and try to put at correct place such that conflict operations order must be maintained same as transactions order.

eg: (T₁ → T₂) — T₂ — T₂ — ...
 ↑ 1 2 3
 T₁ — law of (T₁), check for conflict.



Problems due to Concurrent Executions:

- 1.) WR / Uncommitted read / dirty read
- 2.) R-W / un-repeatable read problem.
- 3.) W-W / lost update
- 4.) Phantom tuple problem

Type of Lock

1. Shared lock : only read
2. Exclusive lock : write / write / read / read / write.

Lock point : position of low level operations / fine-grain lock operations

2PL

- * Ensures Conflict Serializability.
- * Serializing order determined by lock points.
- * 2PL does not ensure recoverability.
- * 2PL suffers from deadlocks ^{may/may not}.
- * 2PL suffers from Starvation.

Rigorous 2PL

- * 2PL + All locks (shared & exclusive) must be held by transaction until commit/rollback.
- * Suffers from Deadlock & Starvation.

Conservative 2PL : Each transaction acquires all the locks in the beginning of execution and releases after commit.

- * Ensures : Conflict serializability, recoverability, cascading, strict-recoverable, no deadlock.
- * But suffers from Starvation.

Time Stamp Protocol

- * Unique time stamp value assigned to each transaction when they arrive in the system.
- * Based on time stamp determine serializability order.
- * Ensures : Conflict Serializability, free from deadlock but not recoverable & Starvation.

2 Phase Locking Protocol

- Growing phase (Acquire locks)
- Shrinking phase (Release locks)

Strict 2PL

- * 2PL + All exclusive locks taken by the transactions must be held until Commit/rollback.
- * Ensures : Conflict serializable, recoverable schedule, Cascading, Strict-recoverable.
- * But suffers from Deadlock and Starvation.

Thomas Write Rule

- * Modified version of time stamping protocol.
- * When T_i attempts to write data item Q , if $T_S(T_i) < W-T_S(Q)$ then T_i is attempting to write an absolute value of $\{Q\}$.
- Rather than rolling back T_i at the Tsp.
- * Allows new Serializable schedules that are not conflict serializable.

Strict Time Stamp = Tsp + Strict schedule

Ordering protocol

→ Ensure: Serializability, recoverability, Cascading Strict Recovery.

RDBMS → Arity: no of attributes (degree)

Relations

Cardinality: no of tuples/records

Table

Row . Columns

(tuple/record) (field/attribute)

Functional Dependency

X, Y — attribute sets of relation R

t_1, t_2 — any two tuples

\Rightarrow If $t_1.x = t_2.x$ then $t_1.y = t_2.y$

must be same

Primal FD (Always valid)

$x \rightarrow y, x \geq y$

Reflexive Rule: $x \geq y$ and $x \rightarrow y$.

Non Primal FD

$x \rightarrow y, x \cap y = \emptyset$, and $x \rightarrow y$
true satisfy

Semi Non Primal FD

$x \rightarrow y, x \not\geq y$ and $x \cap y \neq \emptyset$

Augmentation rule: $x \rightarrow y$,

thus $x_2 \rightarrow y_2$

Transitive rule: $\{x \rightarrow y, y \rightarrow z\}$

thus, $x \rightarrow z$.

Decomp setwise: $\{x \rightarrow y_2\}$, then $x \rightarrow y_1$,
 $x \rightarrow y_2$ and $x \rightarrow z_1$, then $x \rightarrow z_2$

Union: $\{x \rightarrow y, x \rightarrow z\}$, then $x \rightarrow y_1$,
 $x \rightarrow z_1$

Super key: If 'x' is a attribute set of a relation R, then if $[x]^+$ determines all attributes of R, then x is superkey

Candidate key: minimal of Superkey is called candidate key.

Prime attributes: attribute present in any candidate key (not all keys)

→ If X attribute → key / prime attribute, then multiple candidate keys can be present.

Equality b/w 2 FD sets

$F[\dots], G[\dots]$

They are equal if $F \text{ cover } G$, and $G \text{ cover } F$

$F \text{ cover } G$: F covers all FDs of G FD set.

$G \text{ cover } F$: G covers all FDs of F FD set.

→ Maximum no. of candidate keys = $nC_{\lceil \frac{n}{2} \rceil}$ no. of attributes

→ If every single attribute is candidate key, max superkeys = $2^n - 1$

Lossless Join Decompositions

If $R_1 \bowtie R_2 \bowtie R_3 \dots R_n = R$

Binary Method: $R_1 \bowtie R_2$ is lossless, if

* $R_1 \cup R_2 = R$

* Common attribute of R_1 and R_2 's either Superkey of R_1/R_2

Finding minimal cover

(1) Specify the FD such that RHS contains single attribute.

(2) Find redundant attribute on LHS and delete them.

(3) Find redundant FD and delete from FD set.

* Minimal cover can be ≥ 1

Dependency Preserving Decompositions

$$dP \quad F_1 \cup F_2 \cup F_3 \dots \cup F_n = F$$

then dependency preservg.

Normal Forms

1. 1NF
2. 2NF
3. 3NF
4. BCNF

- * Every higher normal form contains lower normal form

Redundancy level: 1NF \rightarrow 2NF \rightarrow 3NF \rightarrow BCNF

First Normal Form (1NF)

- * Relations doesn't contain any multi valued attribute.
- * By default RDBMS is 1NF.

Partial Dependency

$X \rightarrow Y$ is partial FD

if, $A \in Y$ and

$(X - A) \rightarrow Y$.

Second Normal Form (2NF)

- * Every non attribute in R is fully functional dependent on Primary key of R
- * Proper subset of candidate key \nrightarrow non key attribute

Third Normal Form (3NF)

- * Every $X \rightarrow Y$ non-trivial FD must satisfy:

$X \rightarrow Y$,

$\Rightarrow X$: superkey $\sqsubseteq Y$: prime attribute

BCNF

- * Every non-trivial FD must satisfy: $X \rightarrow Y$.
 X : Superkey.
- * If 2 attributes in a relation then always - BCNF

	1NF	2NF	3NF	BCNF
0-1. Redundancy	x	x	x	✓ (suffers from multivalued attributes)
Lossless join	✓	✓	✓	✓
Dependency Preserving	✓	✓	✓	(may/may not)